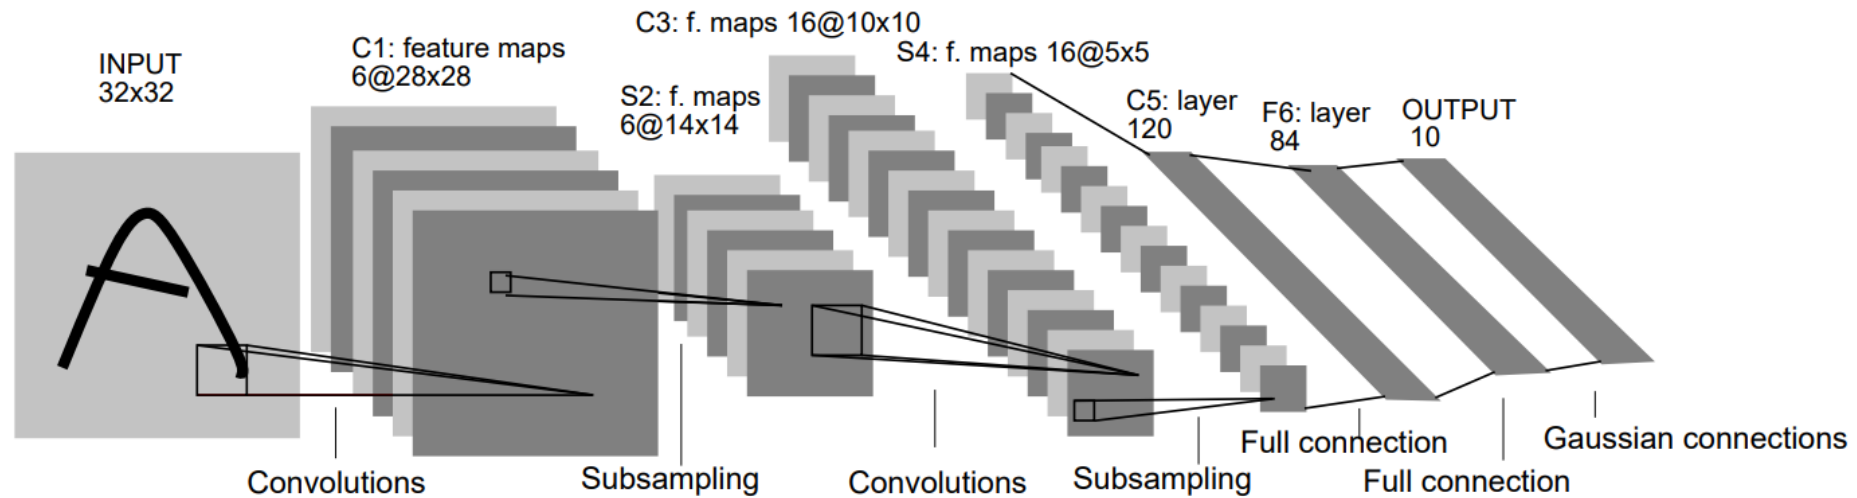


Deep Learning (CS 470, CS 570)

Module 4, Lecture 2: LeNet and AlexNet

LeNet




- LeNet introduced by Yann Lecun in 1998 for Optical Character Recognition (OCR)
- 3 convolutional layers, two pooling layers, and two fully connected layers
- Boundary pixels are ignored during convolution operations therefore no padding is used

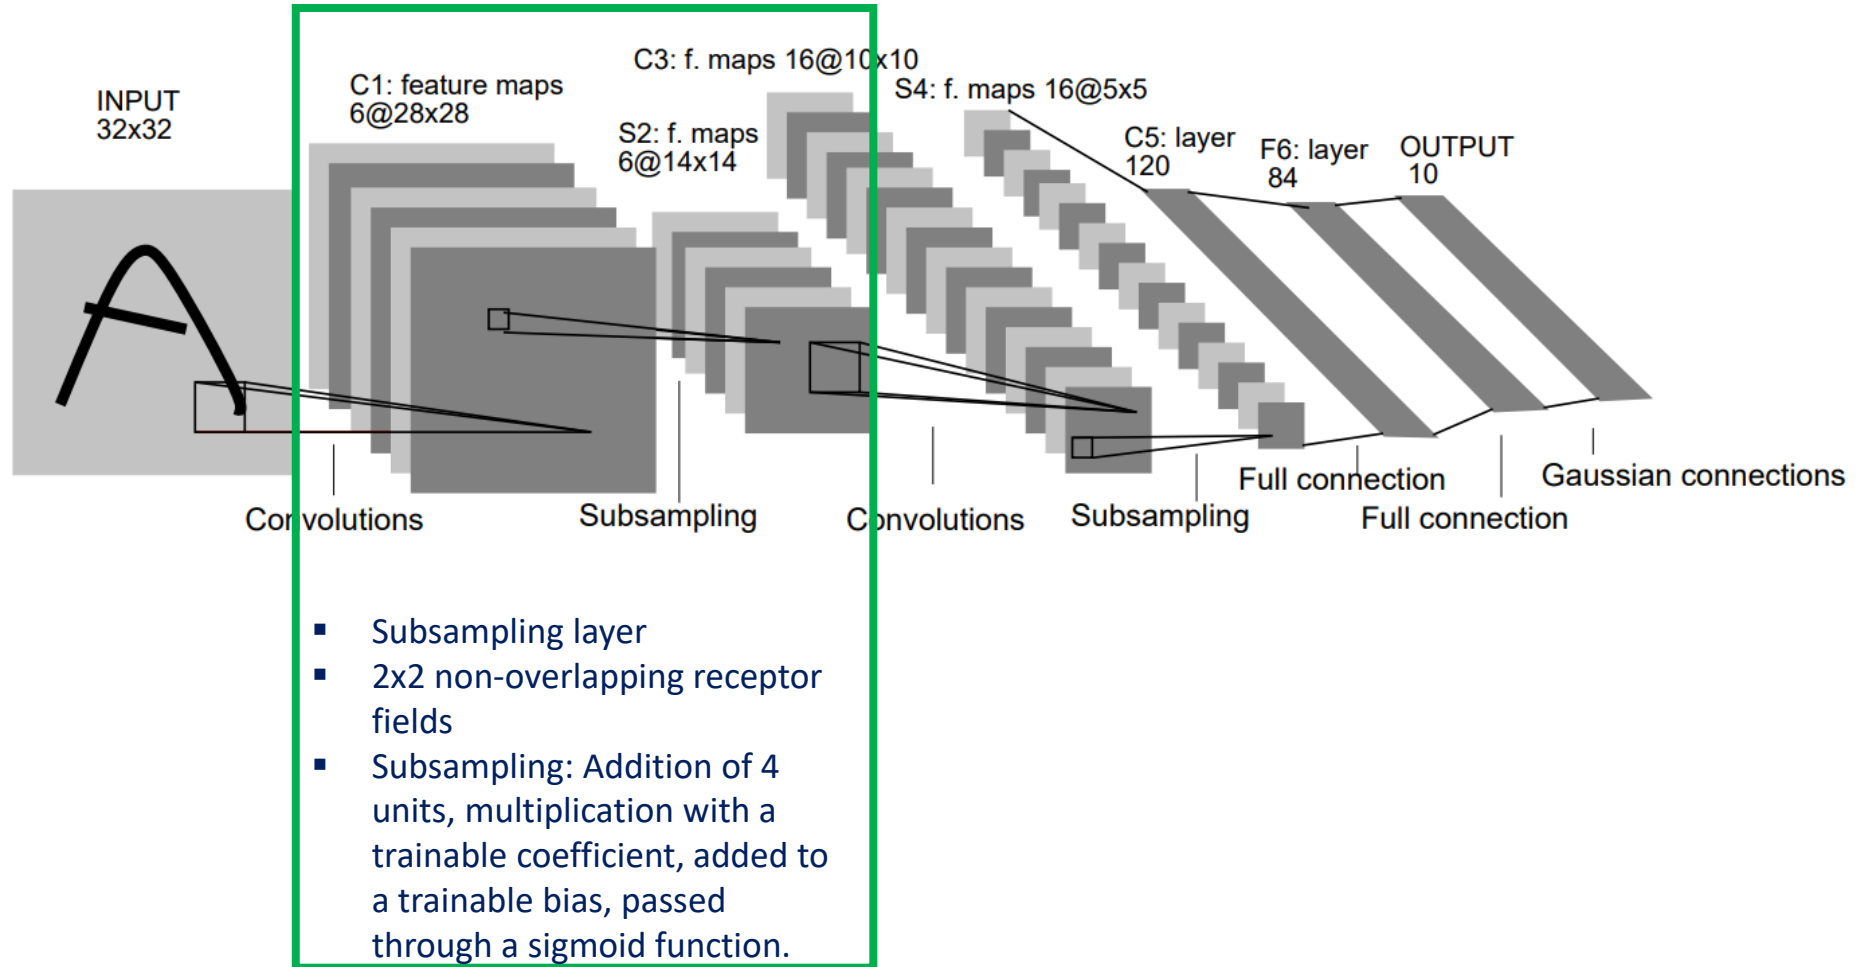
3 6 8 1 7 9 6 6 9 1
 6 7 5 7 8 6 3 4 8 5
 2 1 7 9 7 1 2 8 4 5
 4 8 1 9 0 1 8 8 9 4
 7 6 1 8 6 4 1 5 6 0
 7 5 9 2 6 5 8 1 9 7
 2 2 2 2 2 3 4 4 8 0
 0 2 3 8 0 7 3 8 5 7
 0 1 4 6 4 6 0 2 4 3
 7 1 2 8 9 6 9 8 6 1

Database

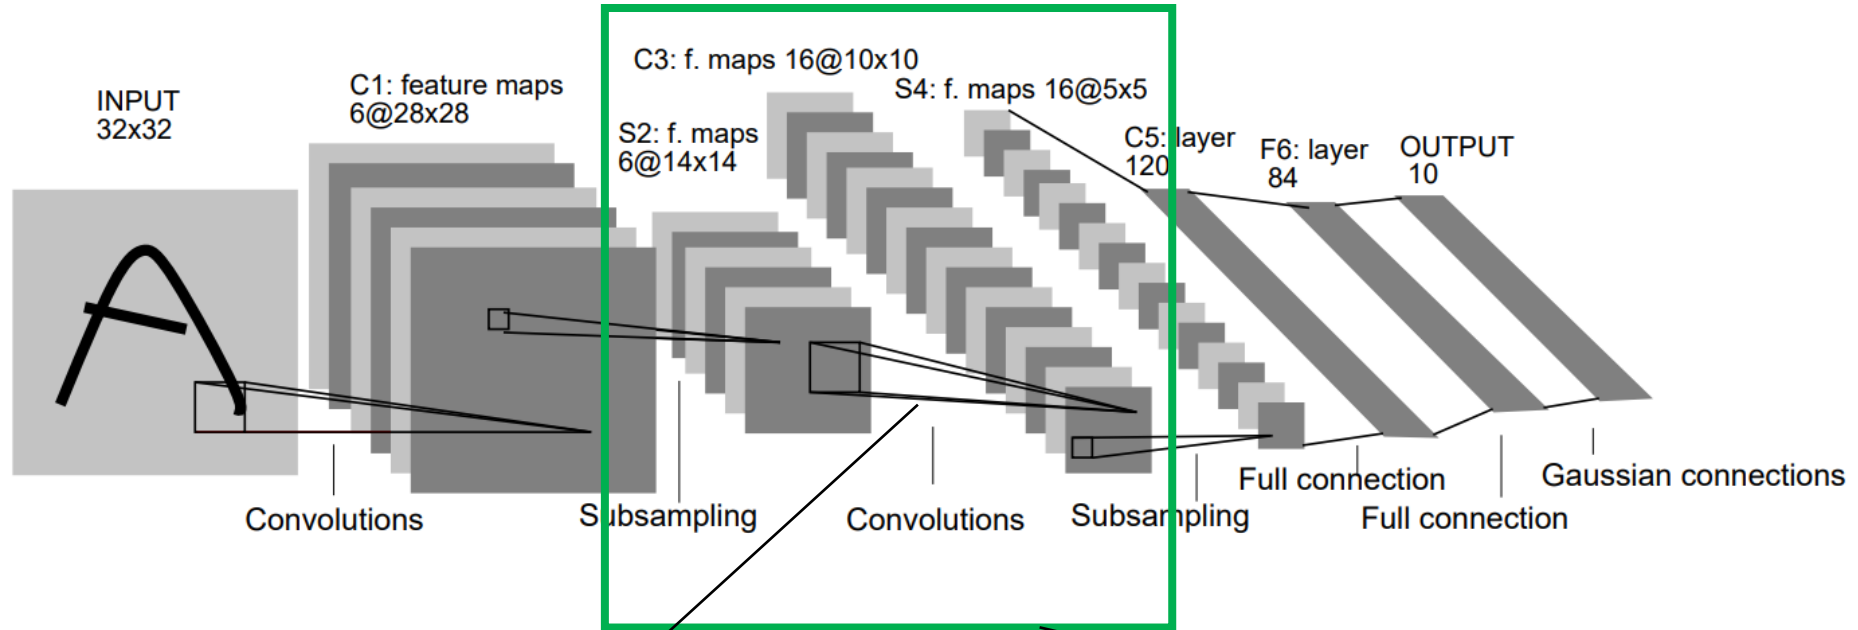
LeNet



LeNet



LeNet



Total trainable parameters: 1516

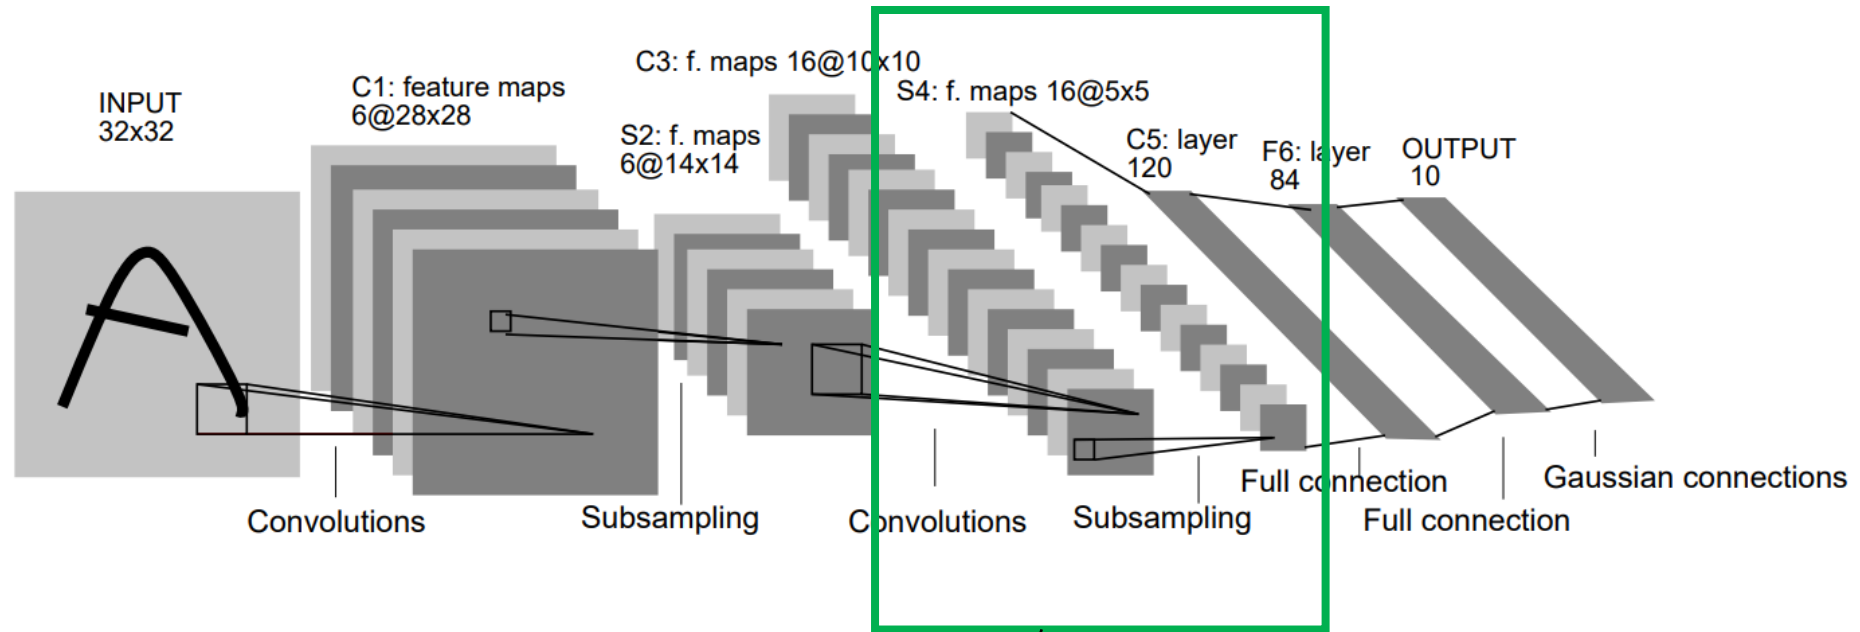
Connection strategy

Input feature maps	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	X				X	X	X			X	X	X	X		X	X
	X	X				X	X	X			X	X	X	X		X
	X	X	X				X	X	X			X		X	X	X
		X	X	X			X	X	X	X			X		X	X
			X	X	X			X	X	X	X		X	X		X

Current layer convolutional filters

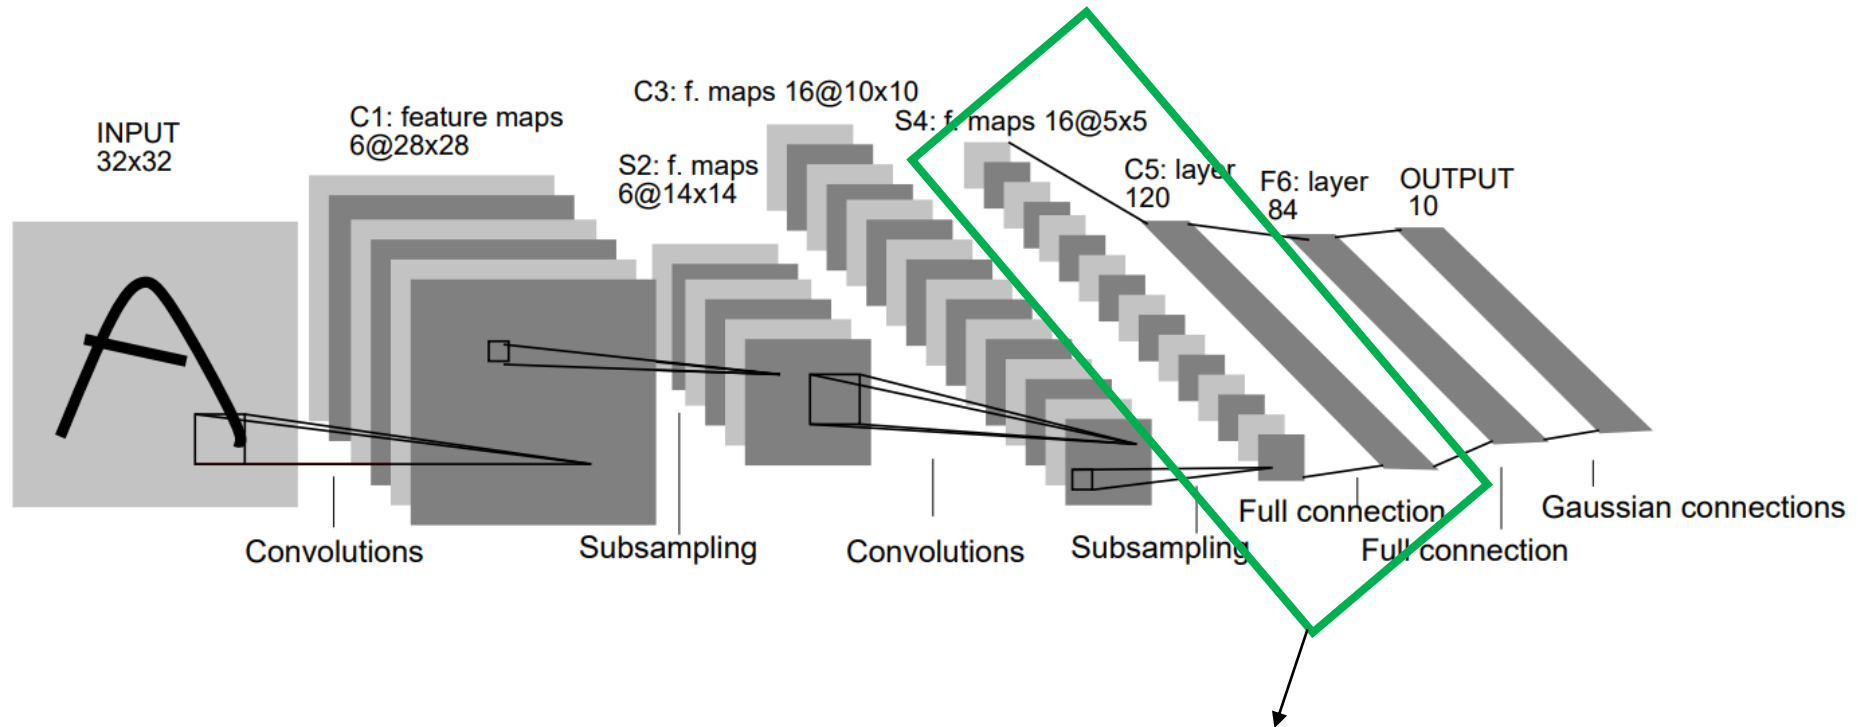
- 16 convolutional filters of size 5x5
- Each of the 16 filters are connected to a subset of input feature maps.
- The reason for subset selection
 - Less parameters to train
 - Different filters are forced to extract different and (hopefully) complementary features

LeNet



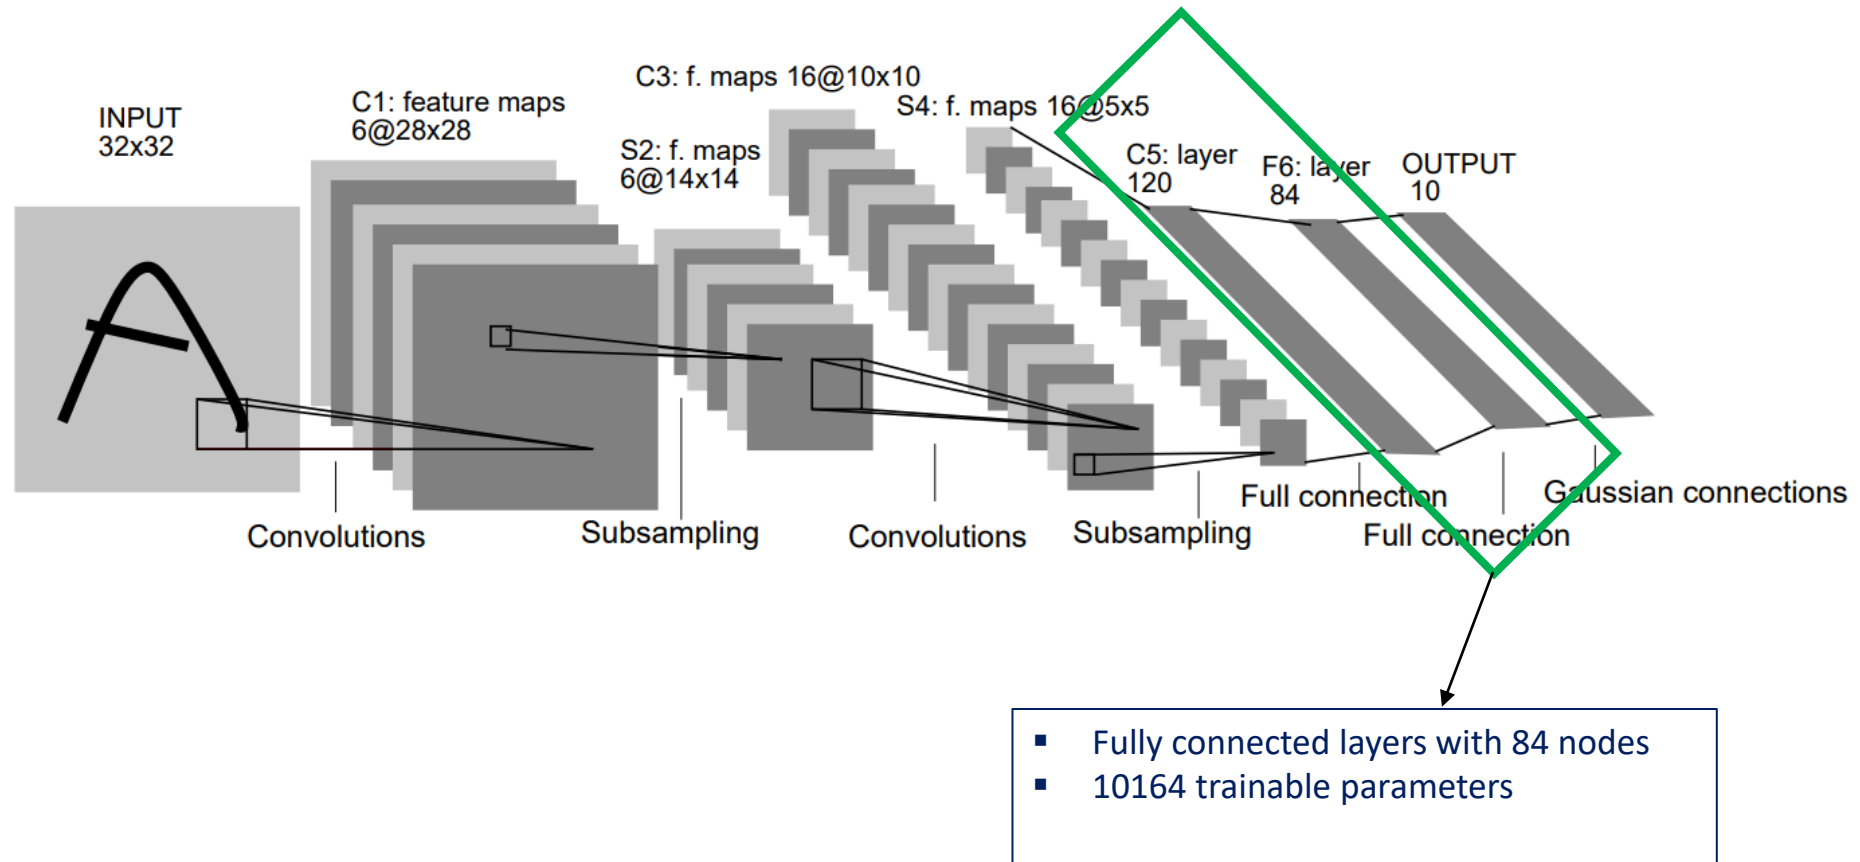
- Subsampling layer
- 2x2 non-overlapping receptor fields
- Subsampling: Addition of 4 units, multiplication with a trainable coefficient, added to a trainable bias, passed through a sigmoid function.

LeNet

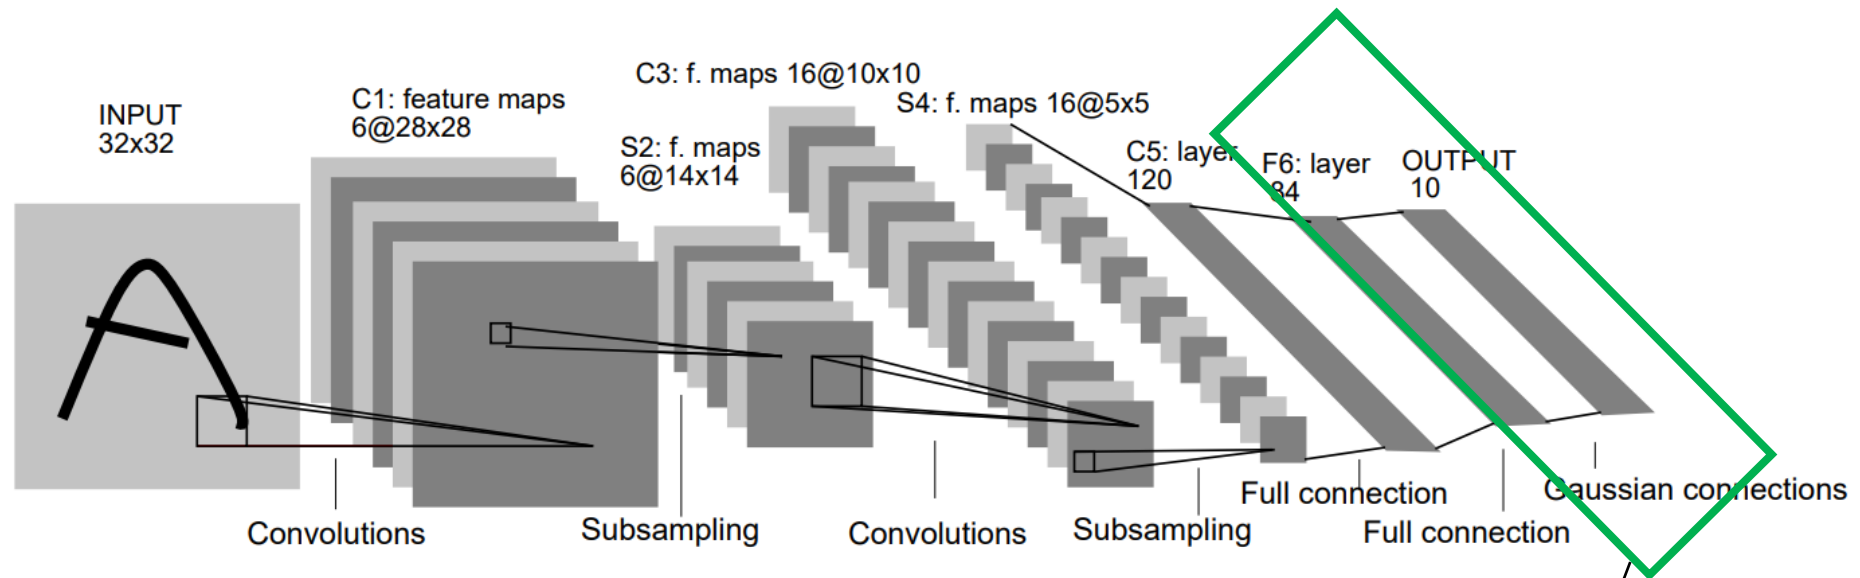


- Fully connected convolutional layer
- Each convolutional filter is connected to all 400 features of previous layer
- 120 convolutional filters of size: 5x5x16
- produce 120 feature maps of size 1x1
- Total trainable parameters: 48120

LeNet

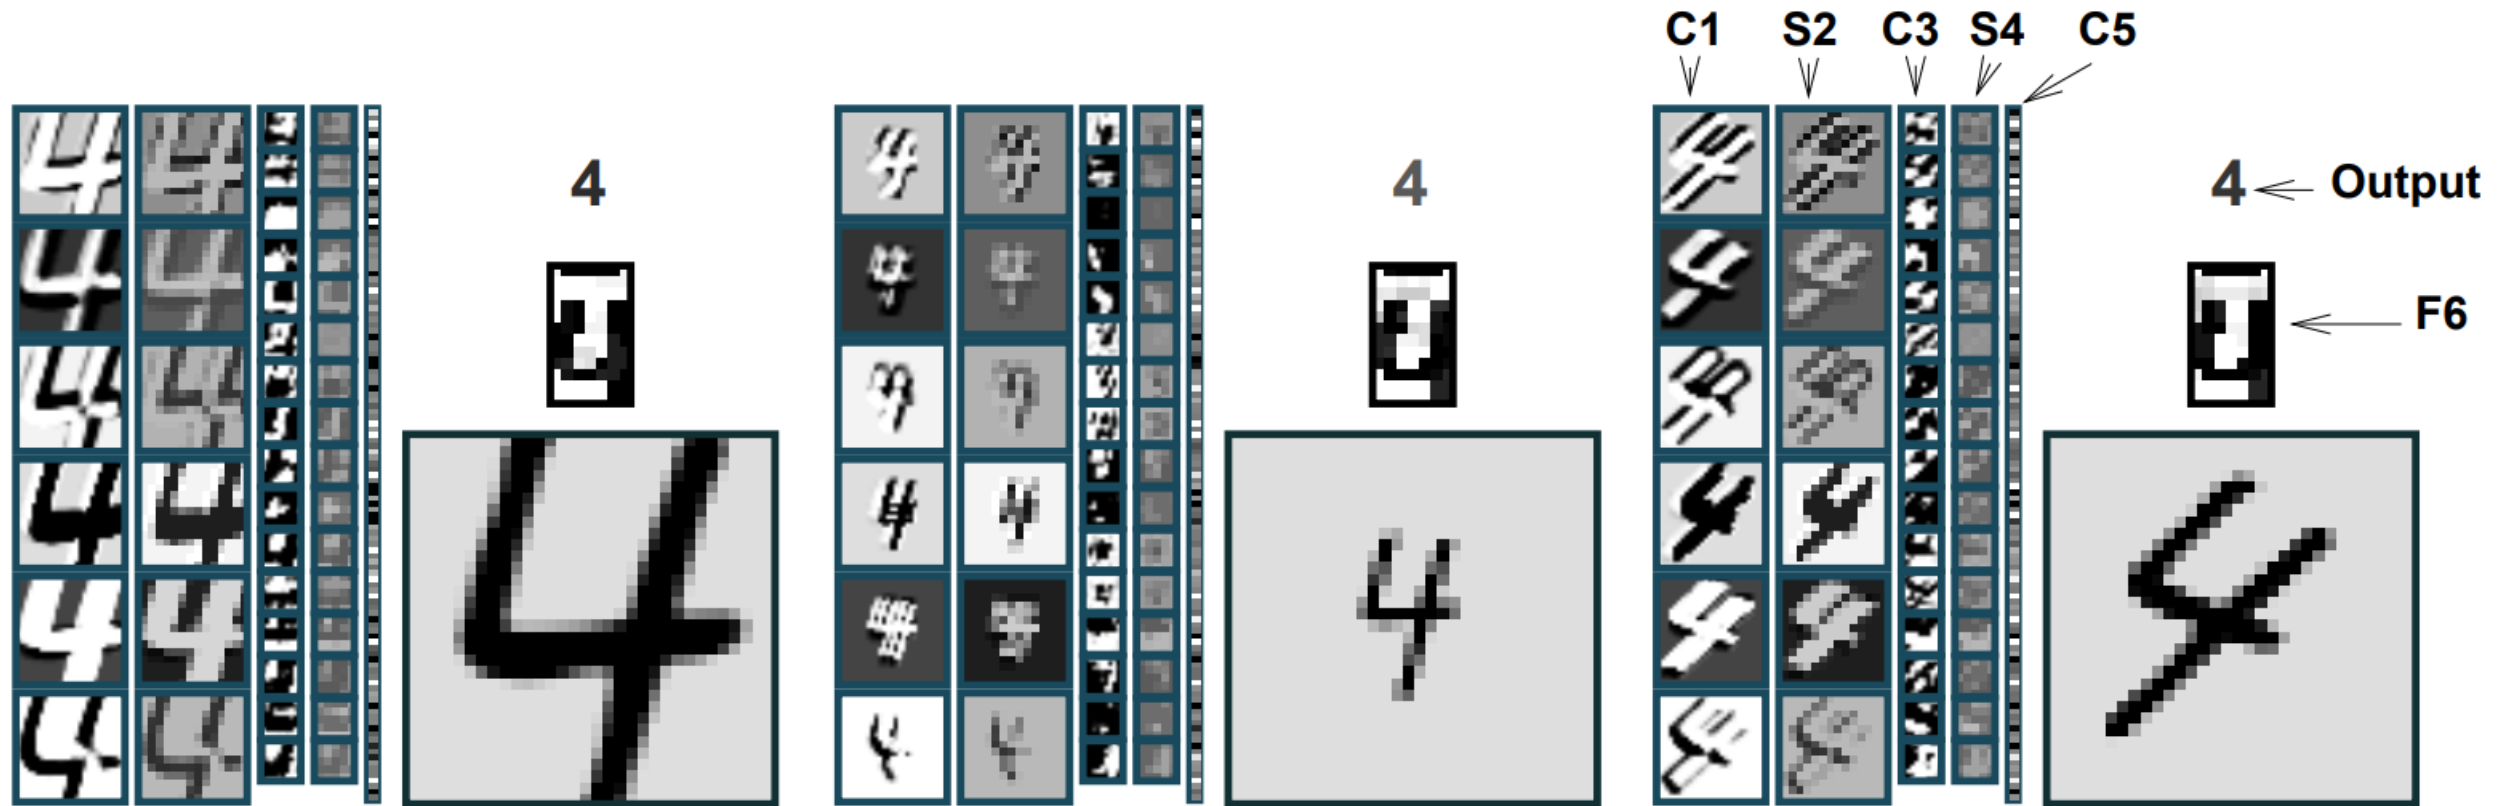


LeNet



- Output layer with 10 nodes
- Connected to each node of previous layer

LeNet: Output



This example shows that LeNet can generate similar responses at F6 for same class (here digit '4') even when there are scale and rotational variations in the input images.

AlexNet

- Even after LeNet, deep learning was not popular because:
 - Lack of computational power didn't allow the network to be deep
 - Shallow architecture was not appropriate for high resolution image classification
- After winning the ImageNet image classification competition, AlexNet became a game changer.
- It is named after Alex Krizhevsky
- First CNN which used GPU for fast computation

ImageNet dataset

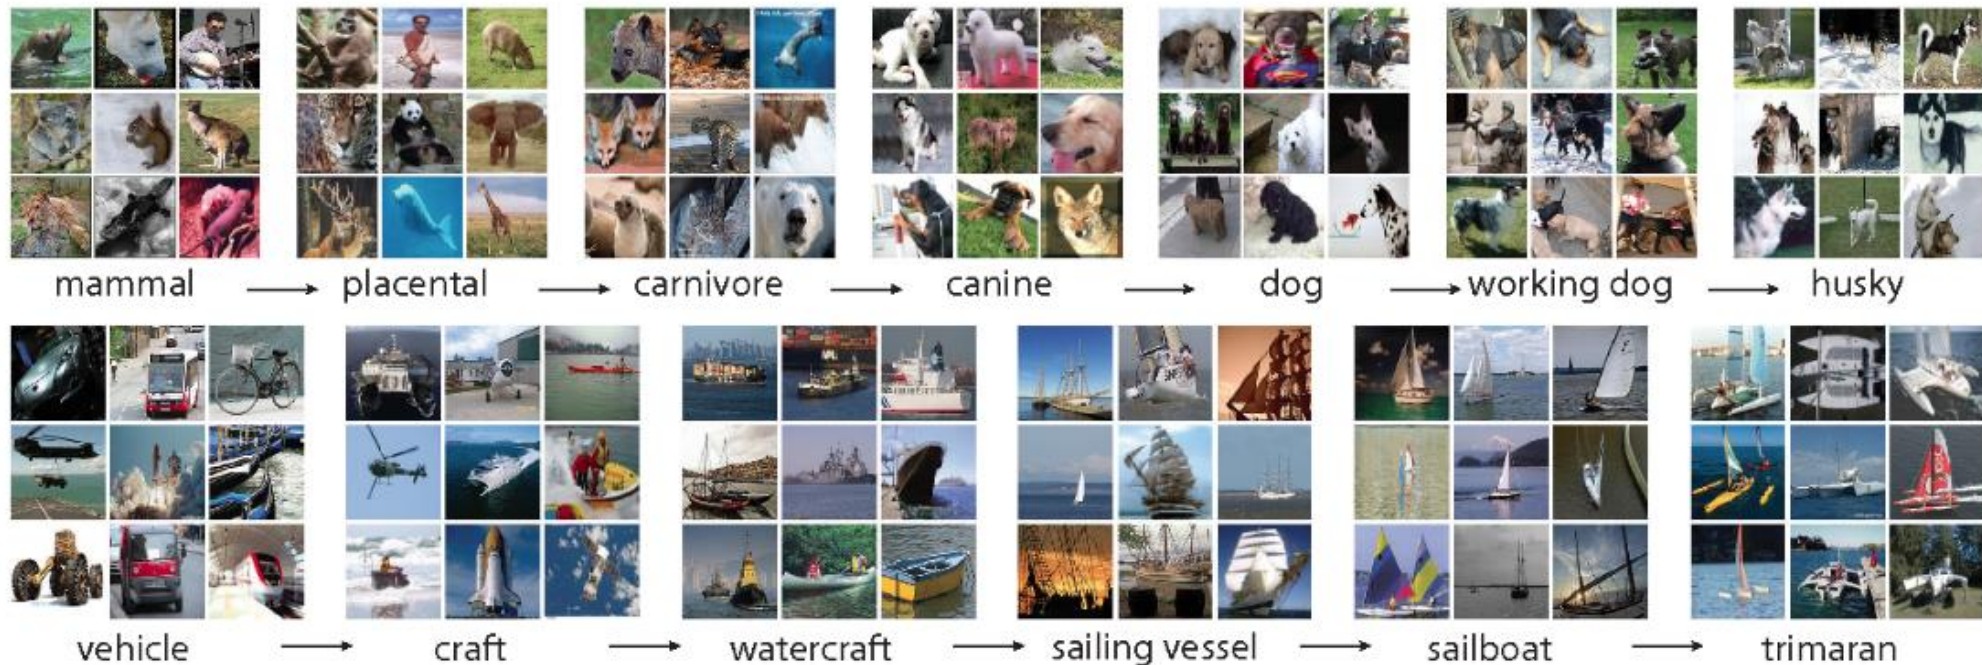



Fig: "ImageNet: A large-scale hierarchical image database", Jia Deng et al.

AlexNet



Database: ImageNet


- 15 million high resolution images
- 22000 different classes
- Manually labeled

<http://www.image-net.org/>

Database used: subset of ImageNet

- 1000 different classes
- Around 1000 images per class
- 1.2 million high resolution training images
- 50,000 validation images
- 150,000 testing images


AlexNet



Network Architecture:


- 8 trainable layers
- 5 convolutional layers, 3 fully connected layers
- Each convolutional layer consist of convolutional filters and a non-linear activation function ReLU
- 3 maxpooling layers for subsampling
- Overall 60 million parameters

AlexNet




- Input images 224x224x3 color images
- 96 convolutional filters of size 11x11x3
- Stride 4, Padding 2
- Output 55x55x96

AlexNet




“ImageNet Classification with Deep Convolutional Neural Networks”, A. Krizhevsky, I. Sutskever, G. E. Hinton

AlexNet




- Input 27x27x96
- 256 filters of size 5x5x48
- Padding 2, Stride 1
- Output: 27x27x256

AlexNet




AlexNet




- Input 13x13x256
- 384 filters of size 3x3x256
- Padding 1, Stride 1
- Output: 13x13x384

AlexNet




AlexNet




- Input 13x13x384
- 256 filters of size 3x3x192
- Padding 1, Stride 1
- Output: 13x13x256

AlexNet




- Input 13x13x256
- Max pooling
- Kernel size 3x3
- Stride 2
- Output: 6x6x256

AlexNet




AlexNet



“ImageNet Classification with Deep Convolutional Neural Networks”, A. Krizhevsky, I. Sutskever, G. E. Hinton


AlexNet



- In total 62.3 million parameters to train
- Conv. Layers: 3.7 million(6%), FC layers: 58.6 million (94%)
- RELU nonlinearity after each convolutional layer

- Input 4096
- Dense connection
- Output: 1000

AlexNet



96 learned filters in convolutional layer 1

Additional Readings

[LeNet](#)

[AlexNet](#)